

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

Searching. Sorting.

Lect. PhD. Arthur Molnar

Babes-Bolyai University

Overview

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search
Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

1 Searching

- The searching problem
- Searching with unordered keys
 - Sequential search
- Searching with ordered keys
 - Binary search
 - Exponential search

2 Sorting

- The sorting problem
- (Binary) Insertion sort
- Quick Sort
- Merge Sort
- TimSort

Let's define the search problem

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

- Data are available in the internal memory, as a sequence of records (k_1, k_2, \dots, k_n)
- We search for the record that has a certain value for one of its fields, called the **search key**.
- We return the position of the record in the given sequence, or -1 if the **search key** was not found
- Depending on the structure of the sequence of records, we distinguish two possibilities:
 - Searching with unordered keys
 - Searching with ordered keys

Searching with unordered keys

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

Problem specification

- **Data:** A is a sequence of records with n elements, where $n \in \mathbb{N}$.
- **Result:** $index$, where $(0 \leq index \leq n - 1)$, the search key's position in sequence A , or $index = -1$, if the key was not found.

Sequential search

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

**Sequential
search**

Searching with
ordered keys

Binary search
Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

Sequential search (iterative and recursive)

Examine the source code in **ex08_sequential_search.py**

What's the computational complexity:

- Best case: search key is the first element checked, so $T(n) = 1 \in \Theta(1)$
- Worst case: search key is not found, so $T(n) = n \in \Theta(n)$

Sequential search

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

**Sequential
search**

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)

Insertion sort

Quick Sort

Merge Sort

TimSort

What's the computational complexity:

- Average case: let's assume¹ 50% of searches end with the element not found, and when the element is found, it has the same probability of being on any position in the list. This time, we simulate $2 * n$ runs; during n runs, the element is not found, and during the remaining runs, it is found on each of the list's n positions;

$$T(n) = \underbrace{\frac{1 + 2 + \dots + n}{2 * n}}_{\text{key found, } n \text{ searches}} + \underbrace{\frac{n + n + \dots + n}{2 * n}}_{\text{key not found, } n \text{ searches}} \in \Theta(n)$$

¹an assumption so that we may provide probability $P(I)$

Sequential search

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

**Sequential
search**

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

Sequential search (iterative and recursive)

Examine the source code in **ex08_sequential_search.py**

What about a recursive implementation²?

$$\begin{cases} T(n) = 1, n = 1 \\ T(n) = T(n-1) + 1, n \geq 1 \end{cases}$$

Leads to a telescopic sum that resolves to $T(n) \in \Theta(n)$

²Only for fun ☺, as the stack usually frowns upon this implementation 🔍🔍🔍

Searching with ordered keys

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

Problem specification

- **Data:** A is a sequence of records with n elements, where $n \in \mathbb{N}$ so that $\forall 0 \leq i < j \leq n - 1, A[i] \leq A[j]$
- **Result:** $index$, where $(0 \leq index \leq n - 1)$, the search key's position in sequence A , or $index = -1$, if the key was not found.

Binary search

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

Binary search (iterative and recursive)

Examine the source code in **ex09_binary_search.py**. Also pay attention to the *test_binary_search* method.

What's the computational complexity:

- Best case: search key is the first element checked, so
$$T(n) = 1 \in \Theta(1)$$
- Worst case: search key is not found.
$$\begin{cases} T(n) = 1, n = 1 \\ T(n) = T(n/2) + 1, n \geq 1 \end{cases}, \text{ each time halving the}$$
array in $O(1)$ time. Since we can do this $\log_2 n$ times, the complexity in the average and worst cases is $\Theta(\log_2 n)$

Exponential search

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching problem
Searching with unordered keys
Sequential search
Searching with ordered keys
Binary search
Exponential search

Sorting

The sorting problem
(Binary)
Insertion sort
Quick Sort
Merge Sort
TimSort

Exponential search

Examine the source code in `ex10_exponential_search.py`. Pay attention to the `binary_search_impl` method.

Determining computational complexity:

We need $\log_2 n$ steps to determine whether the element is in the list; we apply further $\log_2 n$ steps to find its position between indices $i, 2 * i$ (complexity can be determined based on indices i , but for ease of understanding we used n as upper bound, so complexity is $O(\log_2 n)$)

The sorting problem

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

Sorting

Rearrange a data collection in such a way that the elements of the collection verify a given order.

Problem specification

- **Data:** A is a sequence of records with n elements, where $A = (a_1, a_2, \dots, a_n)$, $a_i \in \mathbb{R}$, $i = 1, n$
- **Results:** A' , where A' is a permutation of A , having sorted elements: $a'_1 \leq a'_2 \leq \dots \leq a'_n$.

The sorting problem

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

Terminology:

- **Internal sort** - data to be sorted are available in the internal memory
- **Streaming sort** - data becomes available sequentially (being read from the network, a database etc.)
- **In-place sort** - transforms the input data into the output, only using a small additional space. Its opposite is called out-of-place.
- **Sorting stability** - the original order of multiple records having the same key is preserved
- **Adaptive sort** - an algorithm that takes data structure into account while sorting
- **Hybrid sort** - an algorithm that uses two or more sorting algorithms to speed up sorting

Demo

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)

Insertion sort

Quick Sort

Merge Sort

TimSort

Stable sort example

Examine the source code in **ex12_stable_sort.py**

Sorting algorithms

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

A few algorithms that we will study

- (Binary) Insertion sort
- Quick Sort
- Merge sort
- TimSort

Insertion sort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search
Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

- Traverse the list once, examining each element a single time
- Insert the examined element at the right position in the subsequence of already sorted elements.
- The sub-sequence containing the already processed elements is kept sorted, so that, at the end of the traversal, the entire sequence is sorted.
- Very good for short lists, or lists that are nearly sorted (with only a few elements out of place)

Insertion sort – fact sheet

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

In-place	✓	Stable	✓
Streaming	✓	Space	$O(1)$
Time	Best $O(n)$	Average $O(n^2)$	Worst $O(n^2)$

Binary Insertion sort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

- An optimization of insertion sort
- Since the sequence of already processed elements is sorted, we use binary search to find the place of each new element

Binary Insertion sort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

Question

What happens if use (binary) insertion sort to sort **linked lists** instead of **arrays**? What are the differences?

Binary Insertion sort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

- The searching problem
- Searching with unordered keys
 - Sequential search
- Searching with ordered keys
 - Binary search
 - Exponential search

Sorting

- The sorting problem
- (Binary) Insertion sort
- Quick Sort
- Merge Sort
- TimSort

Example

Examine the source code in **ex13_insertion_sort.py**

Binary Insertion sort – fact sheet

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

In-place	✓	Stable	✓
Streaming	✓	Space	$O(1)$
	Best	Average	Worst
Time	$O(n * \log_2 n)$	$O(n^2)$	$O(n^2)$

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

Quick Sort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

Based on the *divide and conquer* technique

- 1 Divide:** partition array into 2 sub-arrays such that elements in the lower part \leq elements in the higher part.

Partitioning

Re-arrange the elements so that the element called pivot occupies the final position in the sub-sequence. If i is that position: $k_j \leq k_i \leq k_l$, for $Left \leq j < i < l \leq Right$

- 2 Conquer:** recursively sort the 2 sub-arrays.
- 3 Combine:** trivial since sorting is done in place.

Quick Sort - time complexity

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

- An example of in-place, divide & conquer, non-stable sort 😊
- Its time complexity depends on the distribution of splits, namely pivot selection. Some strategies:
 - Choose the first or last element
 - Choose a random element, or the one in the middle
 - Median of the first, middle and last elements in the partition (helps when data is already sorted or sorted in reverse)

Quick Sort - best partitioning

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching problem

Searching with unordered keys

Sequential search

Searching with ordered keys

Binary search

Exponential search

Sorting


The sorting problem

(Binary) Insertion sort

Quick Sort

Merge Sort

TimSort



- We partition n elements $\log_2 n$ times, so
 $T(n) \in \Theta(n \log_2 n)$

Quick Sort - worst partitioning

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

- In the worst case, function Partition splits the array such that one side of the partition has only one element:

$$T(n) = T(1) + T(n-1) + \Theta(n) = T(n-1) + \Theta(n) = \sum_{k=1}^n \Theta(k) \in \Theta(n^2)$$

Quick Sort - Worst case

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting


The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort



- Worst case appears when the input array is sorted or reverse sorted, and we select the first or last element as the pivot
- The n elements are partitioned n times, $T(n) \in \Theta(n^2)$

Quick Sort - Multi-pivot versions

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

- **1961** - The original version of quicksort, using a single pivot is published by Tony Hoare
- **2009** - Research shows that due to modern CPU cache architectures, the constant factor in the algorithm's complexity can be lowered by using 2 pivots

Two-pivot quicksort in Java 8

Implementation by Vladimir Yaroslavskiy (inventor of two-pivot quicksort), Jon Bentley, and Josh Bloch (author of the Java Collections Framework)

<https://hg.openjdk.org/jdk8/jdk8/jdk/file/tip/src/share/classes/java/util/DualPivotQuicksort.java>

Quick Sort - Multi-pivot versions

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

- **2014** - Additional research shows possible further improvements when using 3 pivots

Research paper - Multi-Pivot Quicksort: Theory and Experiments by Kushagra et al.

<https://epubs.siam.org/doi/abs/10.1137/1.9781611973198.6>

Quick Sort – fact sheet

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

In-place	✓	Stable	X
Streaming	X	Space	$O(\log_2 n)$
	Best	Average	Worst
Time	$O(n * \log_2 n)$	$O(n * \log_2 n)$	$O(n^2)$

Sorting

The sorting
problem

(Binary)

Insertion sort

Quick Sort

Merge Sort

TimSort

Merge Sort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)

Insertion sort

Quick Sort

Merge Sort

TimSort

- Based on the divide & conquer approach.
- The list to be sorted is divided in two sub-lists that are sorted separately. The sorted sub-lists are then merged.
- Each sub-list is sorted using the same approach until we get to sub-lists of length 1, which we know are sorted.

Merge Sort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)


Insertion sort

Quick Sort

Merge Sort

TimSort

- + Merge sort has good time complexity and it's easy to write an implementation that works well...
- It has non-constant extra-space complexity
- Time is wasted when merging very short lists (length 1, 2, 4, ...)



Merge Sort – fact sheet

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)

Insertion sort

Quick Sort

Merge Sort

TimSort

In-place	X	Stable	X
Streaming	X	Space	$O(n)$
	Best	Average	Worst
Time	$O(n * \log_2 n)$	$O(n * \log_2 n)$	$O(n * \log_2 n)$

Sorting Algorithm Optimization

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search
Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

- There's no really "best" sorting algorithms, as data size, structure, cost of comparisons, CPU cache sizes and other things influence actual run times
- e.g., merge-insertion sort (a.k.a Ford-Johnson algorithm) was known for a long time as the sorting algorithm with the smallest number of element comparisons
- e.g., pivot count and selection in quicksort implementations
- Combine the strengths of several algorithms into **hybrid algorithms**, which select implementations depending on the size and structure of the input data

Sorting Algorithm Optimization

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

Sorting algorithms visualized

https:

[//www.toptal.com/developers/sorting-algorithms](https://www.toptal.com/developers/sorting-algorithms)

TimSort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)

Insertion sort

Quick Sort

Merge Sort

TimSort

- Developed for Python in 2002 by Tim Peters
- Adopted in Java 7, Android, Swift, V8 and Rust
- A hybrid algorithm built on merge sort that employs binary insertion sort for small sections of the list, leverages natural element ordering and in-cache data 😊

A Word of Warning

You'll find **many** resources on the interwebs showing "TimSort", but most of them are either incomplete, or straightforward combinations of merge and insert sort. Some of them are even buggy 😞

The Real TimSort

- The definitive description of the algorithm by its author (<https://svn.python.org/projects/python/trunk/Objects/listsort.txt>) together with its source code (<https://github.com/python/cpython/blob/main/Objects/listobject.c>)
- Some more details regarding it (<https://mail.python.org/pipermail/python-dev/2002-July/026837.html>)

A few other resources for those interested

- Good (video) descriptions (<https://www.awesomealgorithms.com/home/tim-sort>) and (<https://ericmervin.medium.com/what-is-timsort-76173b49bd16>)
- Detailed, step by step description **with Python source code** (<https://vladris.com/blog/2021/12/30/timsort.html>)
- Source code in Python from the RPython project (<https://github.com/reingart/pypy/blob/master/rpython/rlib/listsort.py>)

TimSort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

- The searching problem
- Searching with unordered keys
 - Sequential search
- Searching with ordered keys
 - Binary search
 - Exponential search

Sorting

- The sorting problem (Binary)
- Insertion sort
- Quick Sort
- Merge Sort
- TimSort**

In the author's own words...

This describes an adaptive, stable, natural mergesort, modestly called timsort (hey, I earned it ☺). It has supernatural performance on many kinds of partially ordered arrays (less than $\lg(N!)$ comparisons needed, and as few as $N-1$), yet as fast as Python's previous highly tuned samplesort hybrid on random arrays.

In a nutshell, the main routine marches over the array once, left to right, alternately identifying the next run, then merging it into the previous runs "intelligently". Everything else is complication for speed, and some hard-won measure of memory efficiency.

<https://svn.python.org/projects/python/trunk/Objects/listsort.txt>

TimSort

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

Most existing lists include sublists that are already sorted (e.g., [1, 4, 5, 2, 99, 6, 9, 11, 10, 90] has [1, 4, 5] and [6, 9, 11] length-3 sublists)

- 1 Determine the value for *minrun*, a parameter that balances good insert sort performance ($32 < \textit{minrun} < 64$) so that the list is divided into a power of 2 runs
- 2 Sort consecutive runs of length *minrun* using binary insertion sort
- 3 Place runs on a stack ($O(n)$ space complexity)
- 4 Merge consecutive runs from the stack using merge sort (this leads to a stable sort)

TimSort - optimizations

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

Sorting

The sorting
problem

(Binary)
Insertion sort

Quick Sort

Merge Sort

TimSort

- While elements are sorted, keep going even after *minrun*
- A run sorted in descending order can be reversed to ascending in $O(n)$ (check order of first two elements)
- Only call insertion sort for those parts of the run that are not already sorted
- Merging often (and soon after runs are completed) raises the chance that data is (still) in the CPU cache
- Introduce a stack invariant on the length of the merges to ensure balanced merges
- Optimize merging by finding which elements remain in the existing order (e.g., best case when the first element of run A is larger than the last of run B)
- When merging uses data from the same run repeatedly, use an exponential search (galloping) to identify how many more elements can be used

TimSort – fact sheet

Lecture 03

Lect. PhD.
Arthur Molnar

Searching

The searching
problem

Searching with
unordered keys

Sequential
search

Searching with
ordered keys

Binary search

Exponential
search

In-place	✓	Stable	✓
Streaming	X	Space	$O(n)$
	Best	Average	Worst
Time	$O(n)$	$O(n * \log_2 n)$	$O(n * \log_2 n)$

Sorting

The sorting
problem

(Binary)

Insertion sort

Quick Sort

Merge Sort

TimSort